


CRITIQUE INTELLIGENCE SYSTEM

A source-linked architecture for turning live or simulated design-review audio into evidence, alternatives, decisions, actions, and revision rationale.



MAP OF VARIATIONS

- A

System spine
end-to-end value and authority
- B

Signal compiler
audio-to-critique extraction contract
- C

Live sequence
fast/slow lanes, latency and recovery
- D

Critique ledger
project memory and artifact revision
- E

Evaluation twin
simulation, faults and measurable gates
- F

Experience & governance
before/during/after surfaces
- G

Production architecture
trust boundaries and scale path
- H


Business wedge
market seam, moat and investor gates

HOW TO READ THIS DOCUMENT

Green is implemented in the repository. Blue is the next product boundary. Amber marks decisions that remain human-authoritative. Red names a risk or evidence gate. Purple marks market logic, not validated demand.


VARIATION A / SYSTEM SPINE


Two clocks, three lanes, one recoverable room





Four surfaces, one explicit authority budget







NEXT SYSTEM

HUMAN AUTHORITY

RISK / GATE

MARKET

VERDICT

FUND DISCOVERY, NOT SCALE

The repo demonstrates a safe extraction spine and a differentiated product hypothesis. It does not demonstrate paid retention, natural-review accuracy, artifact workflow pull, or a durable moat.

Proposed management thresholds

6 / 10

teams complete 4+ reviews

> 50%

eligible recurring reviews captured

> 70%

displayed signals have valid exact anchors

< 5%

false 'decision made' claims

< 20s

median correction time

> 40%

important claims linked to action/issue/revision

3+

paid pilots or priced continuations

Evidence beats narrative

audience beats narrative
teams export summaries and ignore the graph

disposition/revision linking feels like clerical work

usage requires founder or researcher facilitation

reliable extraction arrives too late for live value

security requirements make the first segment uneconomic

buyers compare only with free notes and reject a workflow premium

FOR

meeting capture and artifact review are paid workflows; provenance matters

AGAINST

incumbents already cover notes, tasks, comments, issues, and evidence repositories

UNKNOWN

whether spoken-critique-to-revision is painful enough to buy separately

READ THIS AS

A falsifiable company thesis: teams will repeatedly pay to preserve evidence, alternatives, decisions, and revision rationale that meeting tools and artifact comments leave disconnected. Every gate above can invalidate that thesis.

What exists, what is proposed, what was verified

IMPLEMENTED

NEXT SYSTEM

HUMAN AUTHORITY

RISK / GATE

MARKET

Evidence as of 2026-08-07

UNIT

93 / 93 passing at diagram phase

BROWSER

54 / 54 across Chromium, Firefox, WebKit, phones, tablet

TYPE / BUILD

TypeScript clean; 19-route-entry Next production build

FLY

machine healthy; public Critique Radar inspected

PUBLIC PROBE

exact quotes persisted through provider 4XX and deadline

PACKAGE AUDIT

3 high transitive findings remain; Next major migration gate

●

Microsoft Teams Facilitator — live notes, decisions, tasks, in-person mobile

●

Read / Granola / Fathom — commodity feature and price anchors

●

Figma / Filestage / Autodesk — artifact review and issue workflow

1

Dovetail — source traceability as adjacent trust property

Figma 2026 report — collaborative decision quality is increasingly salient

SOURCE GUIDE

Detailed links, pricing caveats, investor reasoning, risks, and go/no-go gates are in 06-product-and-investor-assessment.md. Architecture and novelty background remain in 01-05. The PDF source is generate_critique_intelligence_system.py