

Critique Intelligence

Turn spoken design-review reasoning into confirmed, source-linked project state—not another meeting summary.

VERDICT

Promising workflow wedge; invest only in paid design-partner discovery.

The code proves a usable extraction spine. It does not yet prove natural-review accuracy, repeated use, willingness to pay, retention, or a defensible data advantage.

1 / MARKET OPENING

A real seam, inside a crowded category

- INITIAL CUSTOMER / HYPOTHESIS  
20–200 person product-design teams and agencies running recurring, consequential reviews.
- JOB TO BE DONE  
Preserve why a design changed: evidence, competing positions, confirmed decisions, owners, and the later artifact revision.
- COMPETITIVE FACT  
Teams, Read, Granola, and Fathom already cover notes and actions. Figma, Filestage, Autodesk, and Dovetail own adjacent artifact or evidence workflows.
- OPENING / INFERENCE  
The under-served seam is spoken critique → human-confirmed decision → exact artifact region or issue → observable revision.
- POSITIONING RULE  
Sell the Critique Ledger workflow. Do not compete as a broad AI meeting assistant.

2 / PRODUCT AND PROOF

Useful vertical slice; incomplete business

- VERIFIED  
Live or simulated final turns flow through ingest, a bounded 10-signal compiler, exact-quote validation, safe fallback, SSE, and the Critique Radar.
- VERIFIED  
93 unit tests; 54 browser cases; TypeScript and production build pass; public Fly deployment is healthy.
- OBSERVED  
A public three-turn probe produced exact-quote critique signals in 12.1 s. This is a smoke test—not an accuracy study.
- BUILD  
Human confirmation/disposition, artifact-region or issue links, revision responses, tenant auth, retention/deletion, durable jobs, and production Postgres.
- AUTHORITY  
The model may propose structure. People alone confirm decisions, meaning, ownership, and consequence. Never rank participants.

3 / COMMERCIAL TEST

Prices are hypotheses; gates decide

- HYPOTHESIS  
\$5k–\$15k paid implementation pilot; then \$300–\$1,000/month per workspace, adjusted by review volume and integrations.
- 90-DAY PASS GATES  
6/10 teams complete ≥4 natural reviews
- ≥50% of eligible reviews captured
- ≥70% valid exact-source anchors
- <5% false confirmed decisions
- median correction time <20 seconds
- ≥40% important claims linked to an artifact
- ≥3 teams pay to continue

STOP IF

Teams export summaries but ignore the ledger; linking feels clerical; founder facilitation is required; or buyers only pay generic note-taker prices.

90-DAY ALLOCATION

Three coupled workstreams; no autonomy theater

- WEEKS 1–4  
Trust extraction  
Natural-review corpus, gold labels, correction UI, signal and decision quality.
- WEEKS 3–8  
Close one workflow  
Disposition + one Figma/Linear-style link + revision response; 6–10 design partners.
- WEEKS 6–12  
Make operable  
Tenant auth, retention/deletion, durable jobs, Postgres, replay, paid continuations.

INVESTOR DILIGENCE

What is known, what can still kill the thesis

- UPSIDE IF PROVEN  
A proprietary, permissioned critique-to-revision graph can improve retrieval, workflow automation, evaluation, and switching cost.
- PRIMARY RISKS  
Crowded substitutes; transcription/diarization error; surveillance concerns; integration friction; weak repeat use; no evidence of willingness to pay.
- SECURITY / RELEASE GATE  
No tenant auth, production retention controls, or durable state yet. Three high transitive Next.js dependency findings remain; a major upgrade is required.